



UE21CS343BB2

Topics in Deep Learning

Dr. Shylaja S S

Director of Cloud Computing & Big Data (CCBD), Centre
for Data Sciences & Applied Machine Learning (CDSAML)

Department of Computer Science and Engineering

shylaja.sharath@pes.edu


**Ack:Anirudh Chandrasekar,
Teaching Assistant**

Topics in Deep Learning


RNN - Truncated Backpropagation Through Time



- We have seen how to perform Backpropagation Through Time in our previous lecture.
- We go n timesteps forward in the forward propagation and then perform backpropagation through all the n timesteps.
- But what if the input sequence is very long causing n to be large.



- But what if the input sequence is very long causing n to be large.
- It becomes expensive to calculate the backpropagation in terms of space since in order to backpropagate we need to save the activations in every timestep.
- This could lead to us running out of physical space.




- This is where Truncated Backpropagation Through Time comes in.

Truncated BPTT defines two constants k_1 and k_2 such that

$$1 \leq k_1 \leq k_2 \leq n$$

where n is the sequence length.



Truncated Backpropagation Through Time

- BPTT(n, n) : This is nothing but the traditional BPTT where we move n timesteps forward followed by n timesteps back.
- BPTT($1, n$): In this approach we move 1 timestep in the forward pass followed by n timesteps back.
 - Example: At $t=1$ we move 1 time step forward and then n timesteps back, at $t=2$ we move another time step forward and n timesteps back(here 2 as we backpropagate all the way to the first time step) and so on.
 - So we move forward through n time steps in total and backpropagate $n(n+1)$ i.e n^2 times.


- BPTT(1,n): In this approach we move 1 timestep in the forward pass followed by n timesteps back.
 - The advantage of this approach is that there is a lot of updating taking place in the early gradients i.e in the early timesteps we are updating the weight matrices and over emphasising on the early timesteps.

- BPTT(k,k): In this approach we move k timestep in the forward pass followed by k timesteps back.
 - Example: At $t=1$ we move k time step forward and then k timesteps back, at $t=2$ we move another k timesteps forward and k timesteps back and till the end of the input sequence.
 - The last timestep need not be (k,k) and this can be used when n is not a multiple of k as well.
 - The advantage of this approach is that we need to save only k activations as compared to all the activations in the previous two approaches.

- BPTT(k,2k): In this approach we move k timestep in the forward pass followed by 2k timesteps back.
 - In BPTT(k,k) we are not backpropagating over certain timesteps at all.
 - To avoid such a scenario BPTT(k,2k) can be used to ensure all the timesteps are backpropagated over.

Note: PyTorch implements only BPTT(n,n) and BPTT(k,k). Any other combination of k_1 and k_2 can be implemented manually.

Truncated Backpropagation Through Time



(a) BPTT

(b) Truncated BPTT

- Use Data As-Is

If the number of timesteps in each sequence is small, such as tens or a few hundred, you can utilize the input sequences as-is. TBPTT has been suggested to have practical limitations of 200 to 400 timesteps. You can reshape the sequence observations as timesteps for the input data if your sequence data is smaller than or equal to this range.

For example, if you had a collection of 100 univariate sequences with 25 timesteps, you could reshape it into 100 samples, 25 timesteps, and 1 feature, or $[100, 25, 1]$.

- Naive Data Split

If your input sequences are large, such as hundreds of timesteps, you may need to divide them up into many contiguous subsequences.

This will need the implementation of a stateful LSTM in Keras, with the internal state being retained across sub-sequence input and only being reset at the conclusion of a true fuller input sequence.

- Naive Data Split

If you had 100 input sequences with 50,000 timesteps, for example, each one might be broken into 100 subsequences with 500 timesteps. One input sequence would provide 100 samples, resulting in a total of 10,000 original samples. Keras' input would have a dimensionality of 10,000 samples, 500 timesteps, and 1 feature, or [10000, 500, 1]. It would be necessary to take care to preserve the state throughout every 100 subsequences and to explicitly or implicitly reset the internal state after every 100 samples.

- Domain-Specific Data Split

Knowing the proper number of timesteps to generate a useful estimate of the error gradient can be difficult.

We can generate a model rapidly using the naïve technique, but the model may not be optimal. Alternatively, while learning the issue, we may utilize domain-specific knowledge to predict the number of timesteps that will be important to the model.

- Systematic Data Split

You can systematically examine a suite of possible subsequence lengths for your sequence prediction challenge rather than guessing at a reasonable number of timesteps.

You might do a grid search over each sub-sequence length and pick the arrangement that produces the best overall model.

Acknowledgements & References

- <https://deeplearning.ai>
- <https://youtu.be/CrMgl3OiNnw?si=pQrpZJUFWMhH-pSK>
- <https://www.codingninjas.com/studio/library/truncated-bptt>



UE21CS343BB2

Topics in Deep Learning

Dr. Shylaja S S

Director of Cloud Computing & Big Data (CCBD), Centre
for Data Sciences & Applied Machine Learning (CDSAML)

Department of Computer Science and Engineering

shylaja.sharath@pes.edu

**Ack:Anirudh Chandrasekar,
Teaching Assistant**